


SVM Homework

Pruthvi Bharadwaj

March 10, 2022


9.3 (a)

```
x1 <- c(3, 2, 4, 1, 2, 4, 4)
x2 <- c(4, 2, 4, 4, 1, 3, 1)
cols <- c("red", "red", "red", "red", "blue", "blue", "blue")
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))
```



(b)

```
#Optimal separating hyperplane
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))
abline(-0.5, 1)
```




(c)

The classification rule here would be:

Classify as Red if $x_2 - x_1 + 0.5 > 0$, **and, classify as Blue** if $x_2 - x_1 + 0.5 < 0$

(d)

```
#Margin for maximal margin hyperplane
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))
abline(-0.5, 1)
abline(-1, 1, lty = 2)
abline(0, 1, lty = 2)
```



(e)

The support vectors for the maximal margin classifier are the points (2,1), (2,2), (4,3) and (4,4)


(f)

Since the 7th observation is not a support vector, a slight change in its position will not affect the maximal margin hyperplane.

(g)


The equation $x_2 = -0.25 + x_1$ will also separate all the observations but is not an optimal hyperplane because the margin is smaller than the optimal option.

```
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))
abline(-0.25, 1)
abline(0, 1, lty = 2)
abline(-0.5, 1, lty = 2)
```



(h)

```
plot(x1, x2, col = cols, xlim = c(0,5), ylim = c(0,5))
abline(-0.5, 1)
abline(-1, 1, lty = 2)
abline(0, 1, lty = 2)
points(5,1, col = "red")
```




9.5 (a)

```
set.seed(9)
x1 <- runif(500) - 0.5
x2 <- runif(500) - 0.5
y <- 1*(x1^2 - x2^2 > 0)
```

(b)

```
plot(x1, x2, col = ifelse(y, "red", "blue"))
```




(c)

```
#Logistic regression
df <- data.frame(x1, x2, y)
fit_glm <- glm(y ~ x1 + x2, data = df, family = binomial)
fit_glm
```

```
##
## Call:  glm(formula = y ~ x1 + x2, family = binomial, data = df)
##
## Coefficients:
## (Intercept)          x1          x2
##   0.05514      0.38587     -0.02653
##
## Degrees of Freedom: 499 Total (i.e. Null);  497 Residual
## Null Deviance:      692.8
## Residual Deviance: 691.2    AIC: 697.2
```

(d)

```
pred_fit <- predict(fit_glm, data.frame(x1,x2))
plot(x1, x2, col = ifelse(pred_fit > 0, "red", "blue"), pch = ifelse(as.integer(pred_fit > 0) == y, 1, 4))
```



In the above plot, the circles are the observations that have been classified correctly and the the crosses are the ones that are misclassified. The decision boundary looks linear.

(e)

```
fit_glm1 <- glm(y ~ poly(x1, 2) + poly(x2, 2), data = df, family = binomial)
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(fit_glm1)
```

```
##
## Call:
## glm(formula = y ~ poly(x1, 2) + poly(x2, 2), family = binomial,
##      data = df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.079e-03 -2.000e-08  2.000e-08  2.000e-08  9.076e-04
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      43.78    3063.63   0.014   0.989
```

```
## poly(x1, 2)1    1360.39  102905.10   0.013   0.989
## poly(x1, 2)2   21374.91  785951.63   0.027   0.978
## poly(x2, 2)1    -119.10   88918.85  -0.001   0.999
## poly(x2, 2)2  -21333.50  788724.67  -0.027   0.978
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6.9276e+02  on 499  degrees of freedom
## Residual deviance: 2.4730e-06  on 495  degrees of freedom
## AIC: 10
##
## Number of Fisher Scoring iterations: 25
```

```
fit_glm2 <- glm(y ~ x1 + x2 + x1*x2, data = df, family = binomial)
summary(fit_glm2)
```

```
##
## Call:
## glm(formula = y ~ x1 + x2 + x1 * x2, family = binomial, data = df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.342  -1.199   1.050   1.144   1.291
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.05206    0.08983   0.580   0.562
## x1           0.38234    0.31036   1.232   0.218
## x2          -0.01969    0.31760  -0.062   0.951
## x1:x2         0.64537    1.12041   0.576   0.565
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 692.76  on 499  degrees of freedom
## Residual deviance: 690.87  on 496  degrees of freedom
## AIC: 698.87
##
## Number of Fisher Scoring iterations: 3
```

```
fit_glm3 <- glm(y ~ x1 + x2 + log(x1) + log(x2), data = df, family = binomial)
```

```
## Warning in log(x1): NaNs produced
```

```
## Warning in log(x2): NaNs produced
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```


```
summary(fit_glm3)
```



```
##
## Call:
## glm(formula = y ~ x1 + x2 + log(x1) + log(x2), family = binomial,
##      data = df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.314e-04 -2.100e-08  2.100e-08  2.100e-08  2.932e-04
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      274.39  276445.15   0.001   0.999
## x1              2225.00 2058197.83   0.001   0.999
## x2             -2678.51 1713550.32  -0.002   0.999
## log(x1)          145.08  387545.42   0.000   1.000
## log(x2)          -49.28  246981.59   0.000   1.000
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1.9107e+02  on 138  degrees of freedom
## Residual deviance: 1.7565e-07  on 134  degrees of freedom
## (361 observations deleted due to missingness)
## AIC: 10
##
## Number of Fisher Scoring iterations: 25
```


(f)

```
pred_fit1 <- predict(fit_glm1, df)
plot(x1, x2, col = ifelse(pred_fit1 > 0, "red", "blue"), pch = ifelse(as.integer(pred_fit1 > 0) == y, 1,
```



(g)

```
#Support Vector Classifier
df$y <- as.factor(df$y)
fit_svc <- svm(y ~ x1 + x2, data = df, kernel = "linear")
pred_svc <- predict(fit_svc, df, type = "response")
plot(x1, x2, col = ifelse(pred_svc != 0, "red", "blue"), pch = ifelse(pred_svc == y, 1, 4))
```




In the above plot, the circles represent observations that have been classified correctly and crosses represent the observations that have been misclassified.

(h)

#SVM with non-linear kernel

```
fit_svm <- svm(y ~ x1 + x2, data = df, kernel = "polynomial", degree = 2)
pred_svm <- predict(fit_svm, df, type = "response")
plot(x1, x2, col = ifelse(pred_svm != 0, "red", "blue"), pch = ifelse(pred_svm == y, 1, 4))
```



- (i) SVM with a polynomial kernel performs better. But logistic regression with non-linear predictors performs the best

9.7 (a)

```
data("Auto")
Auto$Y <- ifelse(Auto$mpg > median(Auto$mpg), 1, 0)
Auto$Y <- as.factor(Auto$Y)
```


(b)

```
set.seed(9)
cost <- data.frame(cost = seq(0.01, 100, length.out = 10))
svm_tune <- tune(svm, Y ~ ., data = Auto, kernel = "linear", ranges = cost)
summary(svm_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 11.12
```

```
##
## - best performance: 0.02301282
##
## - Detailed performance results:
##      cost      error dispersion
## 1    0.01 0.07653846 0.05100638
## 2   11.12 0.02301282 0.01891104
## 3   22.23 0.03583333 0.03245677
## 4   33.34 0.03839744 0.03872235
## 5   44.45 0.03839744 0.03872235
## 6   55.56 0.03839744 0.03872235
## 7   66.67 0.03839744 0.03872235
## 8   77.78 0.03839744 0.03872235
## 9   88.89 0.03839744 0.03872235
## 10 100.00 0.03839744 0.03872235
```

```
plot(svm_tune$performances[,c(1,2)], type = "l")
```



cost=11.12 has the best performance.

(c)

```
#Polynomial Kernel
para <- data.frame(cost = seq(0.01, 100, length.out = 5), degree = seq(1, 100, length.out = 5))
```

```
svm_poly <- tune(svm, Y ~ ., data = Auto, kernel = "polynomial", ranges = para)
summary(svm_poly)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
## 75.0025     1
##
## - best performance: 0.02814103
##
## - Detailed performance results:
##   cost degree      error dispersion
## 1    0.0100    1.00 0.54852564 0.01899199
## 2   25.0075    1.00 0.05102564 0.03419231
## 3   50.0050    1.00 0.03326923 0.02434857
## 4   75.0025    1.00 0.02814103 0.01893035
## 5  100.0000    1.00 0.02814103 0.01893035
## 6    0.0100   25.75 0.54852564 0.01899199
## 7   25.0075   25.75 0.54852564 0.01899199
## 8   50.0050   25.75 0.54852564 0.01899199
## 9   75.0025   25.75 0.54852564 0.01899199
## 10 100.0000   25.75 0.54852564 0.01899199
## 11    0.0100   50.50 0.54852564 0.01899199
## 12   25.0075   50.50 0.54852564 0.01899199
## 13   50.0050   50.50 0.54852564 0.01899199
## 14   75.0025   50.50 0.54852564 0.01899199
## 15  100.0000   50.50 0.54852564 0.01899199
## 16    0.0100   75.25 0.54852564 0.01899199
## 17   25.0075   75.25 0.54852564 0.01899199
## 18   50.0050   75.25 0.54852564 0.01899199
## 19   75.0025   75.25 0.54852564 0.01899199
## 20  100.0000   75.25 0.54852564 0.01899199
## 21    0.0100  100.00 0.54852564 0.01899199
## 22   25.0075  100.00 0.54852564 0.01899199
## 23   50.0050  100.00 0.54852564 0.01899199
## 24   75.0025  100.00 0.54852564 0.01899199
## 25  100.0000  100.00 0.54852564 0.01899199
```

Cost of 100 with degree 1 seems to perform the best

#Radial Kernel

```
params <- data.frame(cost=seq(0.01,100,length.out = 5),gamma=seq(0.1,100,length.out = 5))
svm_radial <- tune(svm, Y ~ ., data = Auto, kernel = "radial", ranges = params)
summary(svm_radial)
```

```
##
## Parameter tuning of 'svm':
```

```
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
## 25.0075 0.1
##
## - best performance: 0.02294872
##
## - Detailed performance results:
##   cost gamma error dispersion
## 1 0.0100 0.100 0.20621795 0.08131124
## 2 25.0075 0.100 0.02294872 0.02807826
## 3 50.0050 0.100 0.02807692 0.03059334
## 4 75.0025 0.100 0.03064103 0.02901234
## 5 100.0000 0.100 0.03064103 0.02901234
## 6 0.0100 25.075 0.56858974 0.04679438
## 7 25.0075 25.075 0.54557692 0.06126516
## 8 50.0050 25.075 0.54557692 0.06126516
## 9 75.0025 25.075 0.54557692 0.06126516
## 10 100.0000 25.075 0.54557692 0.06126516
## 11 0.0100 50.050 0.56858974 0.04679438
## 12 25.0075 50.050 0.56089744 0.05867321
## 13 50.0050 50.050 0.56089744 0.05867321
## 14 75.0025 50.050 0.56089744 0.05867321
## 15 100.0000 50.050 0.56089744 0.05867321
## 16 0.0100 75.025 0.56858974 0.04679438
## 17 25.0075 75.025 0.56602564 0.05072705
## 18 50.0050 75.025 0.56602564 0.05072705
## 19 75.0025 75.025 0.56602564 0.05072705
## 20 100.0000 75.025 0.56602564 0.05072705
## 21 0.0100 100.000 0.56858974 0.04679438
## 22 25.0075 100.000 0.56858974 0.04679438
## 23 50.0050 100.000 0.56858974 0.04679438
## 24 75.0025 100.000 0.56858974 0.04679438
## 25 100.0000 100.000 0.56858974 0.04679438
```

Cost of 25 with gamma 0.1 seems to perform the best


(d)

```
linear <- svm(Y ~ ., data = Auto, kernel = "linear", cost = 11.12)
polynomial <- svm(Y ~ ., data = Auto, kernel = "polynomial", cost = 100, degree = 1)
radial <- svm(Y ~ ., data = Auto, kernel = "radial", cost = 25.0075, gamma = 0.1)


pair_plot <- function(a){
  for (name in names(Auto)[!(names(Auto) %in% c("mpg", "Y", "name"))])
    plot(a, Auto, as.formula(paste("mpg~", name, sep = "")))
}

pair_plot(linear)
```


SVM classification plo




SVM classification plo




SVM classification plo



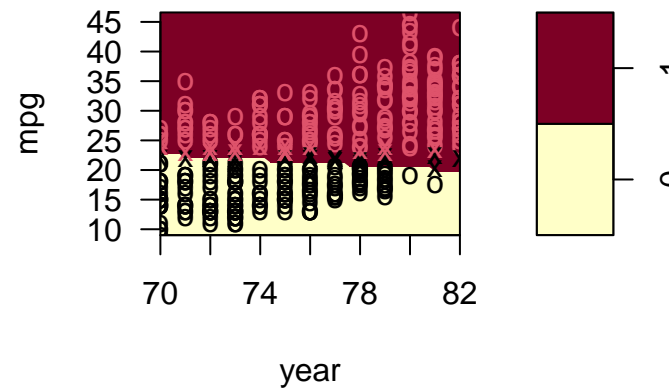
SVM classification plo




SVM classification plo



SVM classification plo



SVM classification plo




plots for linear kernel.

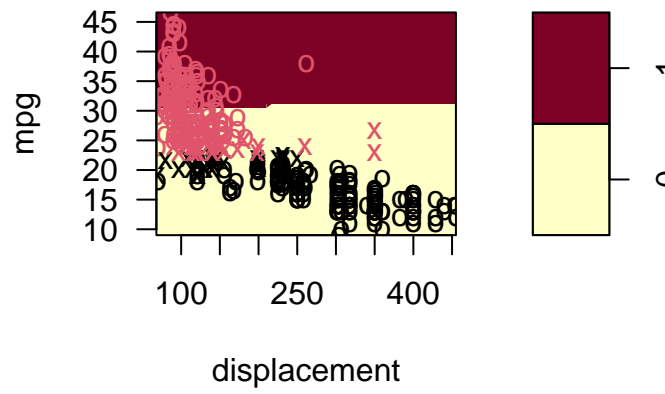
The above are the SVM classification

```
pair_plot(polynomial)
```


SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo



SVM classification plo




SVM classification plo



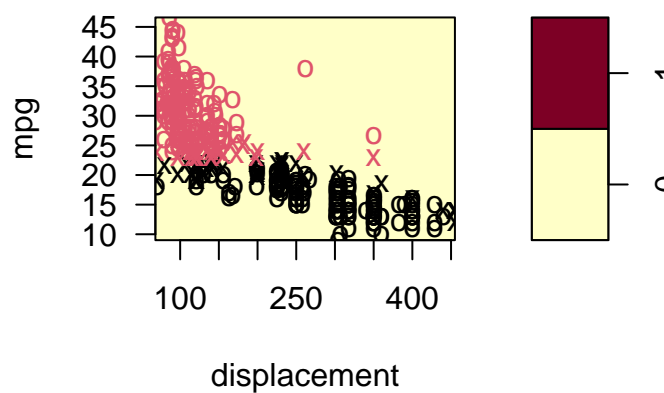
The above plots are the SVM classification plots for polynomial kernel

```
pair_plot(radial)
```


SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo




SVM classification plo



SVM classification plo



SVM classification plo



tion plots for radial kernel

9.8 (a)

```
data("OJ")
set.seed(9)
train_oj <- sample(nrow(OJ), 800)
oj_train <- OJ[train_oj,]
oj_test <- OJ[-train_oj,]
```

(b)

```
oj_svc <- svm(Purchase ~ ., data = oj_train, kernel = "linear", cost = 0.01)
summary(oj_svc)
```

```
##
## Call:
```

The above plots are the SVM classifica-

```
## svm(formula = Purchase ~ ., data = oj_train, kernel = "linear", cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##         cost: 0.01
##
## Number of Support Vectors: 426
##
## ( 213 213 )
##
##
## Number of Classes: 2
##
## Levels:
##  CH MM
```

The SVC creates 432 support vectors out of the 800 training observations. Out of the 432 support vectors, 214 belong to level CH and 213 to level MM.

(c)

#Training error rate

```
pred_train <- predict(oj_svc, oj_train)
table(pred_train, oj_train$Purchase)
```

```
##
## pred_train  CH  MM
##           CH 455  77
##           MM  50 218
```

#Test error rate

```
pred_test <- predict(oj_svc, oj_test)
table(pred_test, oj_test$Purchase)
```

```
##
## pred_test  CH  MM
##           CH 131  39
##           MM  17  83
```

```
(tr_error<- (50+77)/(455+77+50+218))
```

```
## [1] 0.15875
```

```
(te_error <- (39+17)/(131+39+17+83))
```

```
## [1] 0.2074074
```

The training error rate is 15.87% and the test error rate is 20.74%

(d)

```
#For optimal cost

oj_tune <- tune(svm, Purchase ~ ., data = oj_train, kernel = "linear", ranges =
               data.frame(cost = seq(0.01, 10, length.out = 25)))
summary(oj_tune)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   3.34
##
## - best performance: 0.1575
##
## - Detailed performance results:
##       cost  error dispersion
## 1  0.01000 0.16625 0.03175973
## 2  0.42625 0.16125 0.03557562
## 3  0.84250 0.16125 0.03408018
## 4  1.25875 0.16125 0.03197764
## 5  1.67500 0.16125 0.03408018
## 6  2.09125 0.16000 0.03670453
## 7  2.50750 0.15875 0.03586723
## 8  2.92375 0.15875 0.03729108
## 9  3.34000 0.15750 0.03593976
## 10 3.75625 0.15750 0.03593976
## 11 4.17250 0.15750 0.03593976
## 12 4.58875 0.15750 0.03593976
## 13 5.00500 0.15875 0.03586723
## 14 5.42125 0.15875 0.03586723
## 15 5.83750 0.15875 0.03586723
## 16 6.25375 0.15875 0.03586723
## 17 6.67000 0.15875 0.03586723
## 18 7.08625 0.15875 0.03586723
## 19 7.50250 0.15875 0.03586723
## 20 7.91875 0.15875 0.03586723
## 21 8.33500 0.15875 0.03586723
## 22 8.75125 0.15875 0.03586723
## 23 9.16750 0.15875 0.03586723
## 24 9.58375 0.15875 0.03586723
## 25 10.00000 0.15875 0.03586723
```

The optimal cost is 3.75625

(e)

```
#Training error rate
oj_svm <- svm(Purchase ~ ., data = oj_train, kernel = "linear", cost = oj_tune$best.parameters$cost)
svm_train <- predict(oj_svm, oj_train)
table(svm_train, oj_train$Purchase)
```

```
##
## svm_train  CH  MM
##           CH 456 75
##           MM  49 220
```

```
(tr_err <- (49+75)/(456+75+49+220))
```

```
## [1] 0.155
```

```
#Test error rate
svm_test <- predict(oj_svm, oj_test)
table(svm_test, oj_test$Purchase)
```

```
##
## svm_test  CH  MM
##          CH 131 35
##          MM  17 87
```

```
(te_err <- (17+35)/(131+35+17+87))
```

```
## [1] 0.1925926
```

The training error rate is 15.25% and the test error rate is 18.15%.

(f) Radial Kernel

```
oj_radial <- svm(Purchase ~ ., data = oj_train, kernel = "radial")
summary(oj_radial)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj_train, kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##       cost:  1
##
## Number of Support Vectors:  365
##
##   ( 187 178 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
```


The SVM with radial kernel creates 624 support vectors out of the 800 training observations. Out of the 624 support vectors, 313 belong to level CH and 311 to level MM.

```
#Training error rate
```

```
radial_train <- predict(oj_radial, oj_train)
table(radial_train, oj_train$Purchase)
```

```
##
## radial_train  CH  MM
##              CH 466  73
##              MM  39 222
```

```
(tr_err <- (39+73)/(466+73+39+222))
```

```
## [1] 0.14
```

```
#Test error rate
```

```
radial_test <- predict(oj_radial, oj_test)
table(radial_test, oj_test$Purchase)
```

```
##
## radial_test  CH  MM
##              CH 132  37
##              MM  16  85
```

```
(te_err <- (37+16)/(132+37+16+85))
```

```
## [1] 0.1962963
```

The training error rate is 14% and the test error rate is 19.63%.

```
#For optimal cost
```

```
radial_tune <- tune(svm, Purchase ~ ., data = oj_train, kernel = "radial", ranges =
  data.frame(cost = seq(0.01, 10, length.out = 25)))
summary(radial_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 1.25875
##
## - best performance: 0.155
##
```

```
## - Detailed performance results:
##      cost   error dispersion
## 1  0.01000 0.36875 0.04218428
## 2  0.42625 0.16125 0.02461509
## 3  0.84250 0.15500 0.02648375
## 4  1.25875 0.15500 0.03016160
## 5  1.67500 0.15500 0.03238227
## 6  2.09125 0.15625 0.03294039
## 7  2.50750 0.15750 0.03593976
## 8  2.92375 0.16125 0.03557562
## 9  3.34000 0.16000 0.03574602
## 10 3.75625 0.16125 0.03408018
## 11 4.17250 0.16250 0.03535534
## 12 4.58875 0.16625 0.03488573
## 13 5.00500 0.16750 0.03343734
## 14 5.42125 0.16625 0.03682259
## 15 5.83750 0.16625 0.03682259
## 16 6.25375 0.16625 0.03488573
## 17 6.67000 0.16625 0.03488573
## 18 7.08625 0.16625 0.03488573
## 19 7.50250 0.16750 0.03593976
## 20 7.91875 0.16875 0.03448530
## 21 8.33500 0.16875 0.03448530
## 22 8.75125 0.16875 0.03448530
## 23 9.16750 0.17000 0.03343734
## 24 9.58375 0.17000 0.03343734
## 25 10.00000 0.17125 0.03283481
```

The optimal cost ranges from 0.8425 to 1.675

```
#Training error rate
```

```
radial_svm <- svm(Purchase ~ ., data = oj_train, kernel = "radial", cost = radial_tune$best.parameters$cost)
svm_rad <- predict(radial_svm, oj_train)
table(svm_rad, oj_train$Purchase)
```

```
##
## svm_rad  CH  MM
##      CH 469  73
##      MM  36 222
```

```
(tr_err <- (36+73)/(469+73+36+222))
```

```
## [1] 0.13625
```

```
#Test error rate
```

```
svm_rad_test <- predict(radial_svm, oj_test)
table(svm_rad_test, oj_test$Purchase)
```

```
##
## svm_rad_test  CH  MM
##      CH 132  36
##      MM  16 86
```

```
(te_err <- (36+16)/(132+36+16+86))
```

```
## [1] 0.1925926
```

The training error rate is 13.625% and the test error rate is 19.26%

(g) Polynomial Kernel

```
oj_poly <- svm(Purchase ~ ., data = oj_train, kernel = "polynomial", degree = 2)
summary(oj_poly)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj_train, kernel = "polynomial",
##      degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:   1
##   degree:   2
##   coef.0:   0
##
## Number of Support Vectors: 431
##
## ( 219 212 )
##
##
## Number of Classes: 2
##
## Levels:
##  CH MM
```

The SVM with polynomial kernel creates 441 support vectors out of the 800 training observations. Out of the 441 support vectors, 224 belong to level CH and 217 to level MM.

#Training error rate

```
poly_train <- predict(oj_poly, oj_train)
table(poly_train, oj_train$Purchase)
```

```
##
## poly_train  CH  MM
##           CH 468 105
##           MM  37 190
```

#Test error rate

```
poly_test <- predict(oj_poly, oj_test)
table(poly_test, oj_test$Purchase)
```

```
##
## poly_test  CH  MM
##           CH 136 45
##           MM  12 77
```

```
(poly_error<- (105+37)/(468+105+37+190))
```

```
## [1] 0.1775
```

```
(ploye_error <- (45+12)/(136+45+12+77))
```

```
## [1] 0.2111111
```

The training error rate is 17.75% amd the test error rate is 21.11%

```
#For optimal cost
```

```
poly_tune <- tune(svm, Purchase ~ ., data = oj_train, kernel = "polynomial", degree = 2, ranges =
                  data.frame(cost = seq(0.01, 10, length.out = 25)))
summary(poly_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   8.335
##
## - best performance: 0.15875
##
## - Detailed performance results:
##       cost  error dispersion
## 1  0.01000 0.36875 0.04903584
## 2  0.42625 0.20375 0.05684103
## 3  0.84250 0.19375 0.05376453
## 4  1.25875 0.18625 0.04875178
## 5  1.67500 0.18500 0.04556741
## 6  2.09125 0.18375 0.04678927
## 7  2.50750 0.18125 0.04458528
## 8  2.92375 0.18000 0.04338138
## 9  3.34000 0.17375 0.04427267
## 10 3.75625 0.17500 0.04039733
## 11 4.17250 0.17375 0.04226652
## 12 4.58875 0.16875 0.04093101
## 13 5.00500 0.17000 0.04133199
## 14 5.42125 0.16375 0.04619178
## 15 5.83750 0.16375 0.04730589
## 16 6.25375 0.16250 0.04602234
## 17 6.67000 0.16125 0.04505013
## 18 7.08625 0.16250 0.04370037
```

```
## 19  7.50250 0.16250 0.04370037
## 20  7.91875 0.16000 0.04362084
## 21  8.33500 0.15875 0.04489571
## 22  8.75125 0.15875 0.04489571
## 23  9.16750 0.15875 0.04489571
## 24  9.58375 0.16250 0.04714045
## 25 10.00000 0.16375 0.04730589
```

The optimal cost here is between 8.335 to 9.1675

```
#Training error rate
poly_oj <- svm(Purchase ~ ., data = oj_train, kernel = "polynomial", cost = poly_tune$best.parameters$cost)
train_poly <- predict(poly_oj, oj_train)
table(train_poly, oj_train$Purchase)
```

```
##
## train_poly  CH  MM
##           CH 471  80
##           MM  34 215
```

```
(tr_err_poly <- (80+34)/(471+80+34+215))
```

```
## [1] 0.1425
```

```
#Test error rate
test_poly <- predict(poly_oj, oj_test)
table(test_poly, oj_test$Purchase)
```

```
##
## test_poly  CH  MM
##           CH 132  39
##           MM  16  83
```

```
(te_err_poly <- (16+39)/(132+39+16+83))
```

```
## [1] 0.2037037
```

The training error rate is 14.25% and the test error rate is 20.37%

(h)

SVM with linear kernel and with cost 3.75625 gives the best results in terms of the test error rate which is 18.15%